



Microsoft Copilot Studio | Employee Self-Service Agent

ESS Insights Dashboard Interpretation Guide

Get the template and docs at <https://github.com/downesteph/ESS-Insights>

Last updated 7/21/26

The content in this interpretation guide is intended solely to help you better understand the ESS Insights dashboard and the questions it can answer. All figures shown are generated from demo data; once you connect your own Copilot Studio agent, your insights may look materially different.



Introduction

This guide walks through the ESS Insights dashboard page by page and, for each page, explains what the visuals mean and — more importantly — what to do with them. It is written from a consultant's point of view: every metric is paired with an interpretation and a concrete recommendation so you can turn the numbers into action, not just read them.

Detailed Slide Interpretation

For each dashboard page you'll find the page recreated with numbered markers. Each number maps to a callout that poses the key business question, interprets what the data is telling you, and recommends a next step. Follow the numbers in order to read the page the way a consultant would.

Questions?

Reach out to your deployment team or open an issue on the ESS Insights repository.

Frequently Asked Questions

Q: What is the ESS Insights dashboard?

A: A Power BI template that turns Copilot Studio conversation data from your Employee Self-Service agent into adoption, effectiveness, sentiment, and ROI insight — all in one place.

Q: Where does the data come from?

A: From your Copilot Studio agent's conversation transcripts and analytics. The template ships with demo data so you can explore every page before connecting your own agent.

Q: The numbers look different from mine. Why?

A: Every figure in this guide is demo data. Once you connect your own agent, the values, trends, and topic lists will reflect your organization.

Q: How often should I review it?

A: Review leading indicators (adoption, quality, sentiment) weekly and lagging indicators (ROI, business outcomes) monthly. Use the Measuring Value Over Time page to open steering meetings.

Q: What should I do when a metric looks bad?

A: Every detailed page ends with recommendations. Persistent problems flow into the Improvement Opportunities page, which turns weak metrics into a prioritized content backlog with owners.

Glossary

Total Conversations: Count of distinct agent conversations in the selected period.

Total Users: Number of unique people who engaged the agent.

Sessions per User: Average conversations per user; a basic stickiness measure.

Repeat Usage Rate: Share of users who returned for more than one conversation.

DAU / WAU / MAU: Average Daily / Weekly / Monthly Active Users.

Deflection Rate: Share of conversations resolved by the agent without human escalation.

Engaged Rate: Share of conversations where the user actively interacted vs. abandoning.

Abandonment Rate: Share of conversations the user left before reaching a resolution.

Zero-Result Rate: Share of conversations where the agent returned no usable answer.

Turns to Resolve: Average number of back-and-forth exchanges per conversation.

Thumbs-Up Rate: Share of feedback responses that were positive.

CSAT: Customer satisfaction score (1–5) from post-conversation feedback.

Hours Saved / Tickets Deflected: Estimated human effort avoided, based on the assumptions on the Business Impact page.

Leading vs. Lagging Indicators: Leading metrics (quality, adoption, productivity) move first and predict value; lagging metrics (operational KPIs, business outcomes) confirm realized impact.

Detailed Slide Interpretation



1 Why are indicators split into leading and lagging?

INTERPRETATION: Leading indicators (Quality, Adoption, Productivity) move first and predict value; lagging indicators (Operational KPIs, Business Outcomes) confirm the dollars-and-hours impact weeks later.

GUIDANCE: Read the page left-to-right. Early in a rollout expect the left columns to move while the right lag — don't judge ROI in week one.

2 Are the leading indicators healthy?


INTERPRETATION: Quality, Adoption and Productivity are the earliest proof the agent works. They should trend up before you expect business outcomes.

RECOMMENDATION: Set 30/60/90-day targets for each. If adoption or quality stalls, act on Improvement Opportunities before expecting ROI to appear.

3 When can I credibly claim business value?

INTERPRETATION: Operational KPIs and Business Outcomes are the CFO-grade numbers. They only stabilize once adoption is steady and volume is meaningful.

RECOMMENDATION: Report lagging value monthly, not weekly, and pair every ROI figure with the assumptions on the Business Impact page for finance to validate.



4 What belongs here vs. in Copilot Studio?

INTERPRETATION: This is the executive scorecard — the single view for leadership. Operational troubleshooting and content fixes live in Copilot Studio, not on this page.

GUIDANCE: Open every steering meeting with this page; do the actual fix work between meetings in Improvement Opportunities and Copilot Studio.

Organization Adoption

Develop an agent-first experience

1

Is the agent gaining real, repeat users — not one-time curiosity?

INTERPRETATION: 300 users and 779 conversations look healthy, but Repeat Usage (41.7%) and Sessions per User (2.6) are what prove the agent earned a second visit.

RECOMMENDATION: Treat repeat usage as your north-star adoption metric. Below ~30% signals a trust or relevance problem — fix it before scaling licenses.

2

Is adoption growing, plateauing, or decaying?

INTERPRETATION: The weekly trend shows the shape of adoption. A rise-then-plateau is normal after launch; a sustained decline means novelty wore off.


GUIDANCE: Watch the gap between conversations and users. Conversations growing faster than users means your base is getting more value per person.

3

Which departments have adopted — and which haven't?

INTERPRETATION: Users by Organization exposes strength and whitespace. A near-empty department is usually an enablement gap, not disinterest.

RECOMMENDATION: Target bottom-quartile departments with a champion, a use-case demo and a nudge campaign. Adoption spreads fastest team-by-team.



4

Is usage global or concentrated?

INTERPRETATION: Users by Country reveals geographic reach and any language or timezone barriers. Balanced bars suggest the agent generalizes well.

GUIDANCE: If one region lags disproportionately, check language coverage and localized content before assuming there is simply low demand.

Time to Knowledge

How fast users reach an answer

1 How quickly does the agent get people to an answer?

INTERPRETATION: Avg Response Time (13.4s), Duration (57s) and low turns (1.6) mean users get what they need fast — the core promise of self-service.

RECOMMENDATION: Benchmark against your old channel. The consultant story is 'X minutes → under a minute,' and that delta justifies the investment.

2 Where are users giving up?


INTERPRETATION: Abandonment (17.7%), Unengaged (20.4%) and Zero-Result (3.7%) are your friction signals — users who hit a wall and left without an answer.

RECOMMENDATION: Set ceilings (e.g. abandonment <15%, zero-result <5%). Anything above feeds straight into the Improvement Opportunities backlog.

3 Which topics cost the most effort?

INTERPRETATION: Time & Effort by Primary Topic ranks topics by duration and turns. High turns plus long duration means complex or poorly-supported content.

GUIDANCE: Sort by effort and fix the top three topics first. Small content improvements there move the whole-page averages.



4 Does faster always mean better?

INTERPRETATION: Metrics by Session Outcome split speed by result. Fast-but-failed conversations are a red flag — the agent ended quickly because it couldn't help.

RECOMMENDATION: Always read speed together with outcome. Optimize for fast and resolved, never speed alone.

Conversation Outcomes

What actually happens in a conversation

1 How often does the agent fully resolve the request?

INTERPRETATION: Deflection (58.2%) and Engaged Rate (79.6%) are the headline effectiveness numbers — the share handled without a human. Deflection converts directly to cost savings.

RECOMMENDATION: Anchor your ROI story to deflection. A few points of improvement is usually the biggest single lever on the Business Impact page.

2 What is the mix of outcomes?


INTERPRETATION: Agent Session Outcomes shows resolved vs. escalated vs. unresolved. A large escalated slice means the agent is triaging, not solving.

GUIDANCE: Track the resolved slice week over week. Shifts in the donut are your earliest warning that content or model quality changed.

3 Is quality holding as volume grows?

INTERPRETATION: Conversations by Week & Outcome shows whether success holds as usage scales. Rising volume with rising failures means scaling a problem.

RECOMMENDATION: If failures climb with volume, pause expansion and fix content before onboarding more users.



4 Which topics deflect well, and which are hardest?

INTERPRETATION: The Topic Outcomes table plus the most-popular, top-deflected and hardest lists tell you exactly where the agent wins and struggles.

RECOMMENDATION: Promote top-deflected topics in comms to drive volume to proven wins; route the hardest topics into the backlog with named owners.

Agent Feedback

Do users trust the answers?

1

Do users actually trust the answers?

INTERPRETATION: Thumbs-Up Rate (66.4%) and CSAT (3.6/5) are your sentiment signals. The base — 405 responses from 288 users — tells you how representative they are.

RECOMMENDATION: Watch coverage as well as score. A great score from a tiny sample isn't proof; prompt for feedback to make the number trustworthy.

2

Is satisfaction improving as you iterate?

INTERPRETATION: The Weekly Feedback & Satisfaction trend ties sentiment to your release cadence. It should climb after each content or prompt improvement.


GUIDANCE: Annotate this trend with your change log. A dip after a change is a fast, cheap signal to roll back or fix.

3

Which topics delight, and which frustrate?

INTERPRETATION: Thumbs-Up % by Top Topic pinpoints low sentiment even where volume is high. A popular topic with low thumbs-up is a high-impact fix.

RECOMMENDATION: Cross-reference low-sentiment topics with the hardest topics on Conversation Outcomes — overlaps are your top-priority content work.



4

What are users telling you in their own words?

INTERPRETATION: Verbatim Feedback Comments are the qualitative gold — they explain the 'why' behind the scores that numbers alone can't.

GUIDANCE: Read the negatives weekly and theme them. Recurring phrases (wrong answer, couldn't find, too slow) map directly to specific fixes.

Business Impact

What the agent is worth

1

What is the agent worth in hours and dollars?

INTERPRETATION: Tickets Deflected (453), Hours Saved (75.5) and Cost translate usage into business value — the numbers leadership remembers.

RECOMMENDATION: Report net value (hours-saved value minus credit cost), not gross. That's the honest figure that survives finance scrutiny.

2

Are the ROI assumptions defensible?

INTERPRETATION: Every value figure depends on the assumptions shown — Median Hourly Rate (\$20) and Minutes per Ticket (10). Change them and ROI changes proportionally.

RECOMMENDATION: Replace the demo assumptions with your organization's agreed numbers before sharing externally, and get finance to sign off once.

3

Where is value concentrated?

INTERPRETATION: The Estimated Agent Impact Summary breaks value down so you can see which topics or agents drive the savings.

GUIDANCE: Double down on the high-value areas and question the low-value, high-cost ones — not every use case earns its credits.



4

Is spend proportional to value?

INTERPRETATION: The Credit Consumption Leaderboard shows where credits go. High consumption is fine if it maps to high deflection; high spend with low deflection is waste.

RECOMMENDATION: Review top consumers monthly against their deflection and rebalance any agent that burns credits without returning value.

Improvement Opportunities

Turning weak metrics into a backlog

1

Where is the agent failing users right now?

INTERPRETATION: Unrecognized Intent (30.9%), Unanswered (67.3%) and Distinct Unmet Questions (31) quantify the gap between what users ask and what the agent can answer.

RECOMMENDATION: Treat these cards as the agenda for your content team. Each point of unanswered you close is a direct gain to deflection and satisfaction.

2

Which topics need content work first?

INTERPRETATION: Answered / Escalated / Unanswered by Topic shows exactly which subjects fail most. A tall unanswered bar is a specific, fixable content gap.


RECOMMENDATION: Rank topics by unanswered volume and assign the top three to owners this sprint. Fixing high-volume gaps moves every headline metric at once.

3

What's the single highest-leverage action?

INTERPRETATION: The Training Opportunity callout surfaces the most impactful fix the data can find — usually a high-volume, high-failure topic.

GUIDANCE: Make this the first item in every improvement cycle. It's the dashboard telling you where one fix buys the most value.



4

What are we missing entirely?

INTERPRETATION: The 'No Topic Detected' table lists real questions the agent couldn't even categorize — pure whitespace, each with a recommended action.

RECOMMENDATION: Work this table row by row. These are net-new content or routing opportunities that expand what the agent can do, not just how well.